


# MOOC 1: Exploratory Data Analysis for Machine Learning - Study Notes

🔗 **View Original Notion Document:** [https://www.notion.so/MOOC-1-Exploratory-Data-Analysis-for-Machine-Learning-Study-Notes-27b2d55efe0c809b97e6d9d5c2c30ef0?source=copy\\_link](https://www.notion.so/MOOC-1-Exploratory-Data-Analysis-for-Machine-Learning-Study-Notes-27b2d55efe0c809b97e6d9d5c2c30ef0?source=copy_link)

## MODULE 1: AI, ML & DEEP LEARNING

### Core Definitions



©STUDYOPEDIA All rights reserved

### ARTIFICIAL INTELLIGENCE (AI)

- Any program that can sense, reason, act and adapt
- Simulation of intelligent behavior in computers

### MACHINE LEARNING (ML)

- **Subset of AI** - Learns from data, improves over time
- Not explicitly programmed, but learns patterns from data

### DEEP LEARNING (DL)

- **Subset of ML** - Uses multi-layered neural networks
- Automatically learns features, no manual feature engineering needed

## AI History Timeline

- **1950s**: Turing Test developed, AI term officially coined
- **1960s-70s: First AI Winter** - Disappointment with early results
- **1980s**: Expert Systems boom → **Second AI Winter** follows
- **1990s-2000s**: ML breakthrough successes (Google PageRank, speech recognition)
- **2012+**: Deep Learning revolution begins (AlexNet breakthrough)

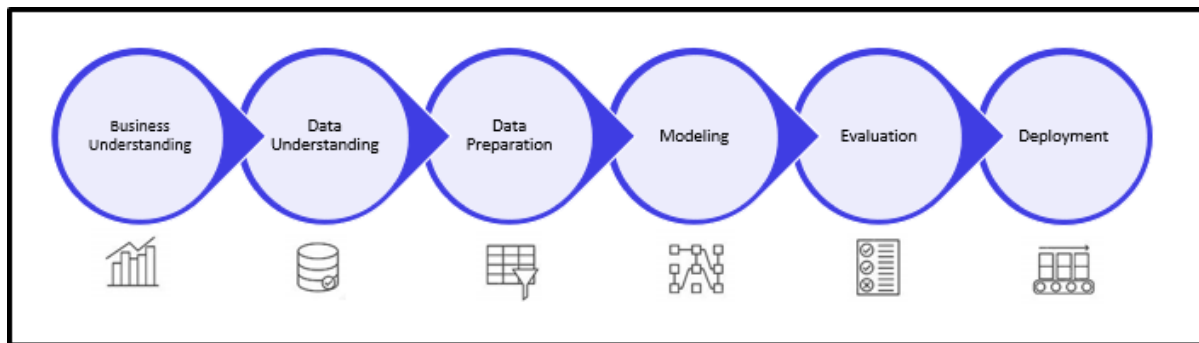
## Real-world Applications

- **Computer Vision**: Face recognition, autonomous vehicles
- **Natural Language Processing**: Machine translation, sentiment analysis
- **Healthcare**: Medical imaging analysis, drug discovery
- **Business Intelligence**: Fraud detection, recommendation engines

## MODULE 2: DATA RETRIEVAL & CLEANING

### Machine Learning Workflow

Problem Statement → Data Collection → Data Exploration & Preprocessing → Modeling → Validation → Deployment



## Detailed Workflow Steps:

1. **Problem Statement:** Define the business problem and success metrics
2. **Data Collection:** Gather relevant data from various sources
3. **Data Exploration & Preprocessing:** Clean, explore, and prepare data
4. **Modeling:** Select and train appropriate ML algorithms
5. **Validation:** Evaluate model performance and generalization
6. **Deployment:** Implement model in production environment

## Data Sources Overview

| Type  | Format             | Python Library                                 | Use Case      |
|-------|--------------------|------------------------------------------------|---------------|
| CSV   | Comma-separated    | <code>pd.read_csv()</code>                     | Tabular data  |
| JSON  | JavaScript Object  | <code>pd.read_json()</code>                    | APIs, NoSQL   |
| SQL   | Structured queries | <code>sqlite3</code> , <code>sqlalchemy</code> | Relational DB |
| NoSQL | Document-based     | <code>pymongo</code>                           | MongoDB       |

## Data Cleaning Strategies

### GARBAGE IN = GARBAGE OUT

Data quality determines model performance!

## Missing Values Handling

- **Remove:** Lose data but no bias introduced
- **Impute:** Replace with mean/median - adds uncertainty
- **Mask:** Treat missing as separate category

# Missing values handling examples

`df.dropna()` # Remove rows with missing values

`df.fillna(df.mean())` # Impute with mean

```
df.fillna(df.median())    # Impute with median
df.fillna('Missing')      # Mask as separate category
```

## Outlier Detection Methods

1. **Visual Methods:** Histograms, Box plots, Scatter plots

2. **Statistical Methods:**

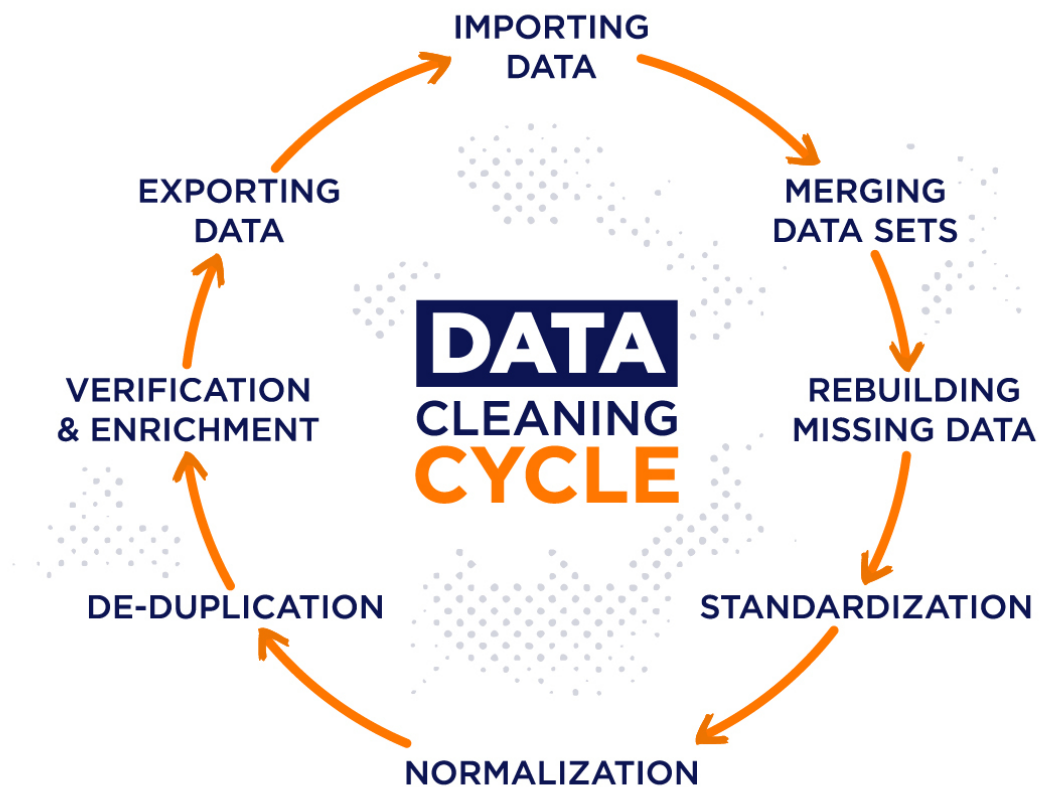
- **IQR Method:**  $Q3 + 1.5 \times IQR$  (upper bound)
- **Z-Score:**  $|z| > 3 \rightarrow \text{outlier}$
- **Residuals:** Standardized, Deleted, Studentized

```
# Outlier detection using IQR method
Q1 = df['column'].quantile(0.25)
Q3 = df['column'].quantile(0.75)
IQR = Q3 - Q1
outliers = df[(df['column'] < (Q1 - 1.5 * IQR)) | (df['column'] > (Q3 + 1.5 * IQR))]

# Z-score method
z_scores = np.abs((df['column'] - df['column'].mean()) / df['column'].std())
outliers = df[z_scores > 3]
```

## Duplicate Data Handling

- **Investigate first:** Distinguish real duplicates vs. valid repeated observations
- **Remove carefully:** Ensure no important information is lost



## MODULE 3: EXPLORATORY DATA ANALYSIS & FEATURE ENGINEERING

### EDA - First Conversation with Data

#### EDA Goals

1. **Understand** data structure & quality
2. **Identify** patterns & relationships
3. **Detect** issues needing cleaning
4. **Guide** modeling decisions

#### Statistical Summary Functions

```
# Core EDA functions  
df.describe()      # Summary statistics
```

```
df.info()          # Data types & nulls
df.corr()          # Correlations
df.value_counts()  # Category frequencies

# Basic visualizations
df['column'].hist() # Histogram
df.boxplot()       # Box plot
sns.pairplot(df)    # Pair plot for relationships
```

## Visualization Libraries

- **Matplotlib:** Foundation library - maximum flexibility
- **Seaborn:** Beautiful statistical plots - `pairplot()`, `heatmap()`
- **Pandas:** Quick plotting - `df.plot()` for rapid insights

## VISUALIZATION TYPES FOR EDA



## Feature Engineering Techniques

### Variable Transformations

- **Log Transform:** Reduces positive skew → normal distribution
- **Polynomial Features:**  $x$ ,  $x^2$ ,  $x^3$  to capture non-linear relationships
- **Box-Cox:** Automatically finds optimal transformation

```
# Variable transformations
import numpy as np
from sklearn.preprocessing import PolynomialFeatures
from scipy.stats import boxcox

# Log transformation
df['log_column'] = np.log1p(df['column']) # log(1+x) to handle zeros

# Polynomial features
poly = PolynomialFeatures(degree=2)
X_poly = poly.fit_transform(X)
```

## Encoding Categorical Variables

### ENCODING METHODS

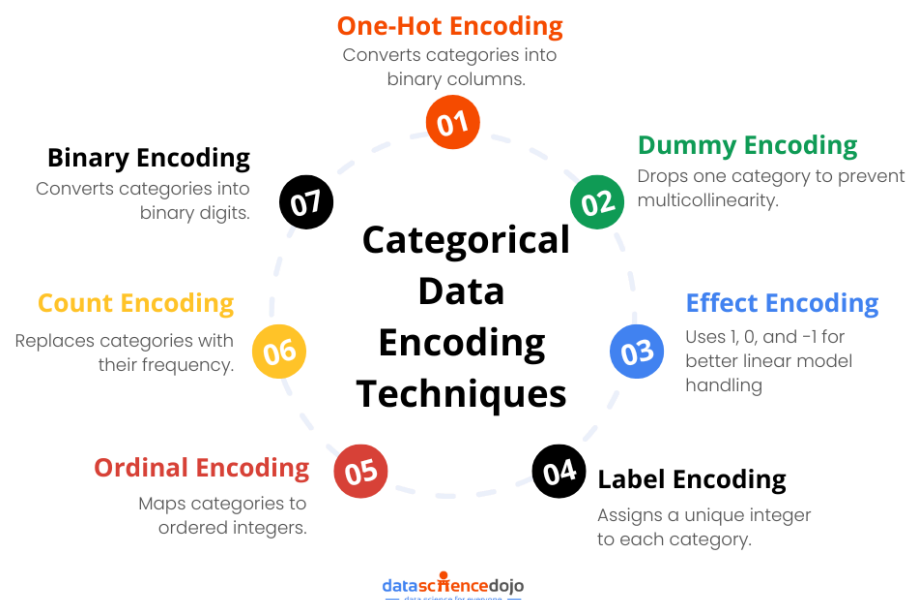
- **Binary Encoding:** True/False → 0/1
- **One-Hot Encoding:** Multiple categories → multiple binary columns
- **Ordinal Encoding:** Low/Medium/High → 1/2/3 (preserves order)

```
# Categorical encoding examples
from sklearn.preprocessing import LabelEncoder, OrdinalEncoder

# One-hot encoding
df_encoded = pd.get_dummies(df, columns=['category_column'])

# Binary/Label encoding
le = LabelEncoder()
df['encoded_column'] = le.fit_transform(df['category_column'])

# Ordinal encoding (preserves order)
ordinal_encoder = OrdinalEncoder(categories=[['Low', 'Medium', 'High']])
df['ordinal_encoded'] = ordinal_encoder.fit_transform(df[['ordered_column']])
```



## Feature Scaling Comparison


| Method           | Formula                   | Range     | Outlier Effect |
|------------------|---------------------------|-----------|----------------|
| Standard Scaling | $(x - \mu)/\sigma$        | Unbounded | Moderate       |
| Min-Max Scaling  | $(x - \min)/(max - \min)$ | [0,1]     | High Risk      |
| Robust Scaling   | $(x - median)/IQR$        | Unbounded | Low Risk       |

```
# Feature scaling examples
from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler

# Standard scaling
scaler = StandardScaler()
X_standard = scaler.fit_transform(X)

# Min-Max scaling
min_max_scaler = MinMaxScaler()
X_minmax = min_max_scaler.fit_transform(X)

# Robust scaling (recommended for outliers)
robust_scaler = RobustScaler()
X_robust = robust_scaler.fit_transform(X)
```



## MODULE 4: STATISTICAL INFERENCE & HYPOTHESIS TESTING

### Key Statistical Concepts



## ESTIMATION vs INFERENCE

- **Estimation:** Point estimate (e.g., mean = 25)
- **Inference:** Interval + uncertainty (e.g., 95% CI: 20-30)

## Parametric vs Non-parametric Approaches

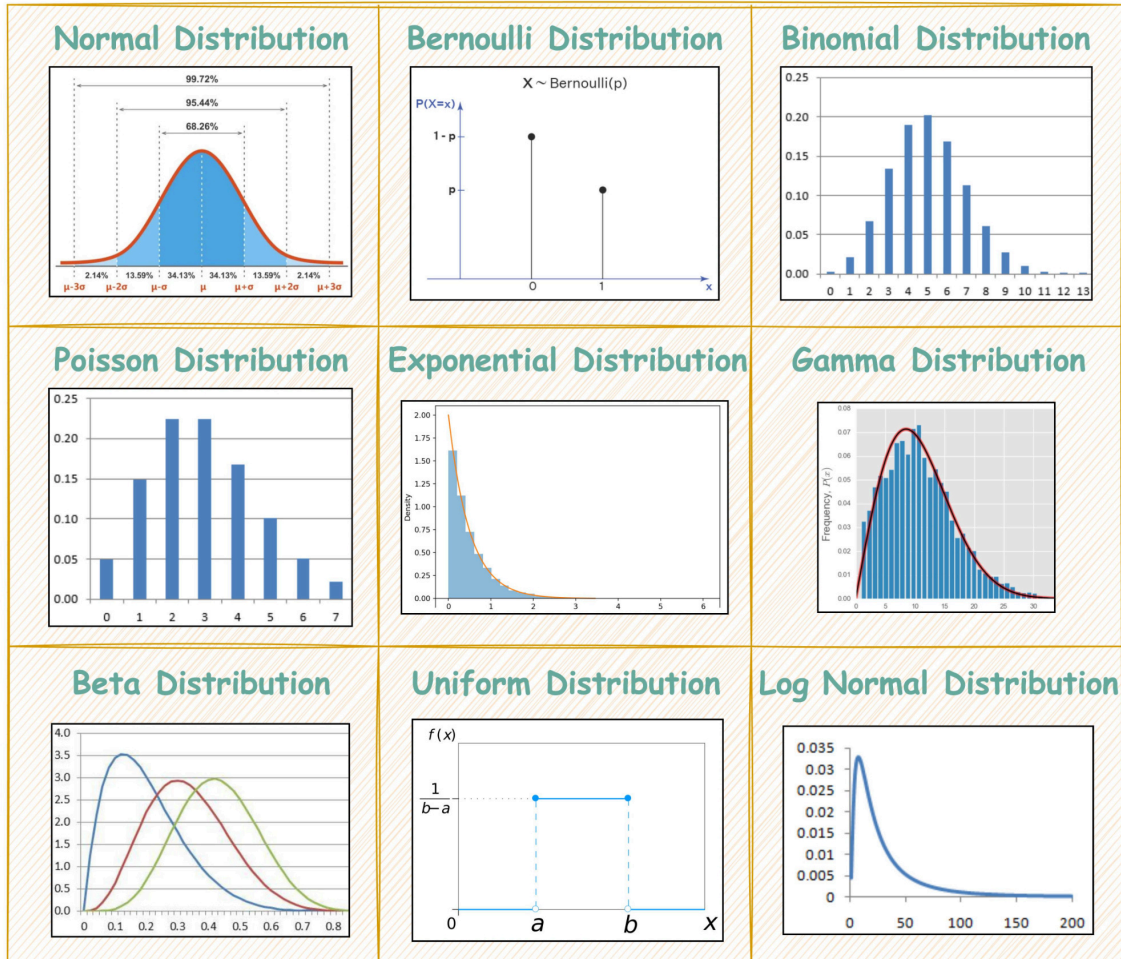
- **Parametric:** Assumes specific distribution (e.g., Normal) + finite parameters
- **Non-parametric:** Distribution-free, relies on actual data structure

## Common Statistical Distributions

| Distribution | Use Case               | Parameters                      | Example                 |
|--------------|------------------------|---------------------------------|-------------------------|
| Normal       | Heights, test scores   | $\mu$ (mean), $\sigma$ (std)    | Most natural phenomena  |
| Binomial     | Success/Failure trials | $n$ (trials), $p$ (probability) | Coin flips, A/B testing |
| Poisson      | Count events in time   | $\lambda$ (rate parameter)      | Website visits per hour |
| Exponential  | Time between events    | $\lambda$ (rate parameter)      | Time to next customer   |



## Most Important Distributions in Data Science



## Hypothesis Testing Framework

### HYPOTHESIS SETUP

- $H_0$  (Null Hypothesis): No effect exists, status quo
- $H_1$  (Alternative Hypothesis): Effect exists, change from status quo
- $\alpha$  (Significance Level): Type I error tolerance (typically 0.05)

### Error Types in Hypothesis Testing

- **Type I Error ( $\alpha$ ):** False Positive - Incorrectly reject true  $H_0$
- **Type II Error ( $\beta$ ):** False Negative - Fail to reject false  $H_0$

- **Statistical Power:**  $1 - \beta$  = Probability of correctly rejecting false  $H_0$

## P-value Interpretation

| P-value = Probability of observing data this extreme if  $H_0$  is true

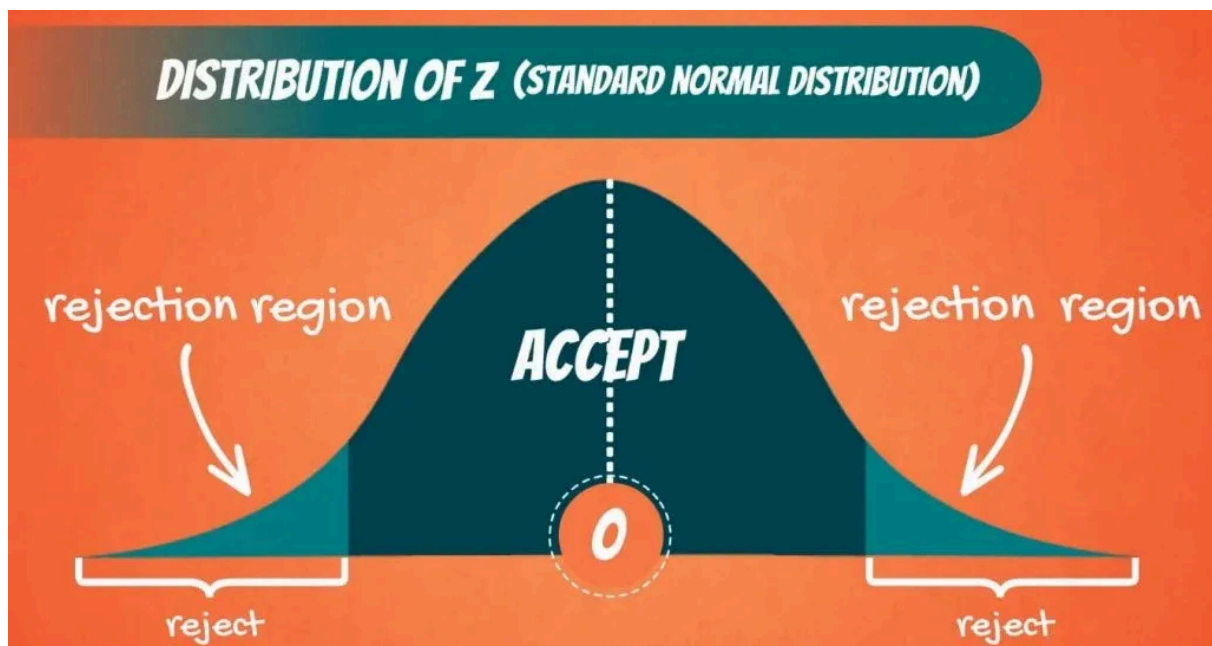
### Decision Rules:

- $p < 0.05$ : Reject  $H_0$  (statistically significant result)
- $p \geq 0.05$ : Fail to reject  $H_0$  (insufficient evidence)

### MULTIPLE TESTING PROBLEM

With  $n$  independent tests:  $P(\geq 1 \text{ Type I error}) \approx n \times \alpha$

**Solution:** Bonferroni Correction  $\rightarrow$  Use  $\alpha/n$  for each test




## Correlation vs Causation

### CORRELATION $\neq$ CAUSATION

#### Four Possible Relationships:

1.  $X \rightarrow Y$ : X directly causes Y
2.  $Y \rightarrow X$ : Y directly causes X (reverse causation)
3.  $Z \rightarrow X, Y$ : Confounding variable Z causes both
4. **Spurious**: Random correlation, no real relationship



### Business Examples

- **Incorrect Interpretation:** More customer service calls → Lower satisfaction
- **Correct Interpretation:** Lower satisfaction → More customer service calls

## BUSINESS APPLICATION: CUSTOMER CHURN PREDICTION

### CHURN PREDICTION WORKFLOW

**Target Variable:** Customer leaves (1) or stays (0)

**Feature Variables:** Tenure, purchase history, demographics, usage patterns

**Model Output:** Probability score (0.0 to 1.0)

**Hypothesis Testing Example:**

- $H_0$ : Customer tenure has no effect on churn probability
- $H_1$ : Longer tenure significantly reduces churn probability

## PYTHON TOOLS & LIBRARIES

### Essential Libraries

```
import pandas as pd      # Data manipulation and analysis
import numpy as np       # Numerical computing and arrays
import matplotlib.pyplot as plt # Basic plotting and visualization
import seaborn as sns    # Statistical data visualization
import scipy.stats as stats # Statistical functions and tests
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

## Key Functions by Category

```
# Data Loading and I/O
pd.read_csv(), pd.read_json(), pd.read_sql()

# Exploratory Data Analysis
df.describe(), df.info(), df.corr()
sns.pairplot(), sns.heatmap(), sns.boxplot()

# Data Preprocessing
StandardScaler(), MinMaxScaler(), LabelEncoder()
pd.get_dummies() # One-hot encoding

# Statistical Testing
scipy.stats.ttest_1samp() # One-sample t-test
scipy.stats.pearsonr()    # Pearson correlation test
scipy.stats.chi2_contingency() # Chi-square test
```